

*****QUESTION 3.1a- using kkn*****

#Set working directory to load the dataset

```
setwd("~/Downloads/IYSE6501/hw2-SP22/data 3.1")
```

```
ccdata <- read.table("credit_card_data-headers.txt",header=TRUE)
```

```
View(ccdata)
```

#Set the seed for reproducibility

```
set.seed(1)
```

#Install kkn package to use kkn function:

```
install.packages("kkn")
```

```
library(kkn)
```

```
set.seed(2)
```

#Train the KNN model using cross-validation

#This following code is reference from office hour on Jan 20:

```
train_model <- train.kkn(as.factor(R1)~., ccdata, kmax=22, scale=TRUE) #I chose kmax=22  
because I simply want to test k from 1 to 22
```

#this line of code is from office hour on Jan 20 to predict the train model:

```
fitted(train_model)[[4]][1:nrow(ccdata)]
```

#Check the best parameters

```
train_model$best.parameters
```

```
> train_model$best.parameters
```

```
$kernel
```

```
[1] "optimal"
```

```
$k
```

```
[1] 12
```

#Split the data to do kkn model based on the k and kernel we found above, 70% for training and 30% for testing:

```
train_set <- sample(1:nrow(ccdata), size=0.7*nrow(ccdata), replace=FALSE)
```

```
train_data <- ccdata[train_set,]
```

```
test_data <- ccdata[-train_set,]
```

#Train the model based on the best parameters found above, the train set and the test set that are split above:

```
model_k12 <- kkn(as.factor(R1)~.,train=train_data,test=test_data, k=12, kernel="optimal",  
scale=TRUE)
```

```
#get prediction from the model with k=12:  
predic_model_k12 <- fitted(model_k12)
```

```
#Calculate test accuracy:  
#This code is referenced from HW1:
```

```
test_accuracy = sum(predic_model_k12==test_data$R1)/nrow(test_data)  
test_accuracy      0.842639593908629
```


```
# Create a plot for misclassification error vs k
```

```
# Misclassification rates for different k values
```

```
misclass_errors <- train_model$MISCLASS
```

```
# this code is referenced from ChatGPT:
```

```
plot(1:22, misclass_errors, type='b', xlab="Number of Neighbors (k)", ylab="Misclassification  
Rate",  
     main="Cross-Validation: Misclassification Rate vs k")
```



As you can see, misclassification is lowest at k=12, so this model works for k=12, kernel="optimal", and the accuracy of 84.26%

*****QUESTION 3.1b*****

#Load data

```
ccdata1 <- read.table("credit_card_data-headers.txt", header=TRUE)
```

#Define factors and response in the dataset

```
response <- as.factor(ccdata1$R1)
```

```
predictors <- ccdata1[,1:10]
```

#install the package kkn

```
install.packages("kkn")
```

```
library(kkn)
```

#Set the seed for the reproduce:

```
set.seed(1)
```

#Split the data, 70% for training:

```
train_data <- sample(1:nrow(ccdata1), size=0.7*nrow(ccdata1),replace=FALSE)
```

```
training_data <- ccdata1[train_data,]
```

```
remaining_data <- ccdata1[-train_data,]
```

#The remaining 30% of data, use half of it for validating and half of it for testing:

```
validate_set <- sample(1:nrow(remaining_data), size=0.5*nrow(remaining_data),  
replace=FALSE)
```

```
validation_data <- remaining_data[validate_set,]
```

```
testing_data <- remaining_data[-validate_set,]
```

#Those loop codes are referenced from HW1, 2-2-3 solution:

#Finding the Best k Using Cross-Validation on the Training Set:

```
check_accuracy = function(X){ #define a function check_accuracy to compute the accuracy for  
different values of k in KNN by iterating over all training data and making predictions.
```

```
  predicted <- rep(0,(nrow(training_data))) # predictions: start with a vector of all zeros
```

```
  # for each row, estimate its response based on the other rows
```

```
  for (i in 1:nrow(training_data)){
```

```
    # data[-i] means we remove row i of the data when finding nearest neighbors...
```

```
    #...otherwise, it'll be its own nearest neighbor!
```

```
    model=kkn(R1~.,training_data[-i,],training_data[i,],k=X, scale = TRUE) # use scaled data
```

```

# record whether the prediction is at least 0.5 (round to one) or less than 0.5 (round to zero)

predicted[i] <- as.integer(fitted(model)+0.5) # round off to 0 or 1
}

# calculate fraction of correct predictions

accuracy = sum(predicted == training_data[,11]) / nrow(training_data)
return(accuracy)
}

# Now call the function for values of k from 1 to 20

acc <- rep(0,20) # set up a vector of 20 zeros to start
for (X in 1:20){
  acc[X] = check_accuracy(X) # test kkn with X neighbors
}
acc

# Find the best k based on validation data.
best_k <- which.max(acc)

# Train knn model with the best k on the training data
best_model <- knn(R1 ~ ., training_data, validation_data, k = best_k, scale = TRUE)

#Evaluating the Model on the Validation Set:
# Predict on validation data, this code is referenced from HW1 solution:
predicted_validation <- as.integer(fitted(best_model) + 0.5)

# Calculate accuracy in validation data:
validation_accuracy <- sum(predicted_validation == validation_data$R1) / nrow(validation_data)
print(paste("Validation accuracy:", validation_accuracy))

```

```

[1] "Validation accuracy: 0.857142857142857"

```

```

# Train final knn model with best k on the combined training and validation data
#This code is referenced from ChatGPT:

```

```
final_model <- kkn(R1 ~ ., rbind(training_data, validation_data), testing_data, k = best_k, scale = TRUE)
```

```
#Evaluating the Model on the Test Set
```

```
# Predict on test data, this code is referenced from HW1 solution:
```

```
predicted_test <- as.integer(fitted(final_model) + 0.5)
```

```
# Calculate test accuracy
```

```
test_accuracy <- sum(predicted_test == testing_data$R1) / nrow(testing_data)
```

```
[1] "Test accuracy: 0.868686868686869"
```

So the accuracy for this model with 70% training data, 15% validation data, 15% test data, without overlapping each other, is 86.86%.

*****Question 3.1a using caret package*****

#Set file as working directory:

```
setwd("~/Downloads/IYSE6501/hw2-SP22/data 3.1")
```

#Load the data in R and call it cc_data:

```
cc_data <- read.table("credit_card_data-headers.txt", header=TRUE)
```

#Install the kernlab package to use kkn function:

```
install.packages("kkn") #for k-nearest-neighbor
```

```
library("kkn")
```

```
install.packages("caret") #for cross validation
```

```
library("caret")
```

#Set the seed for reproduce:

```
set.seed(1)
```

#Indicate the response value is at column R1:

```
cc_data$R1 <- factor(cc_data$R1)
```

```
predictors <- cc_data[, -which(names(cc_data) == "R1")] # Exclude the response variable from the predictors
```

#Start separating data into a training set, and a test set:

#Start on a training set, take 70% the cc_data:

```
training_set <- sample(1:nrow(cc_data), size=0.7*nrow(cc_data), replace=FALSE)
```

#replace=FALSE make sure there is no overlapping data

#To avoid overlapping, I do an extra step to avoid overlapping, and indicate the train set and the remaining set:

```
train_set <- cc_data[training_set,]
```

```
test_set <- cc_data[-training_set,]
```

#Code reference: <https://www.geeksforgeeks.org/svm-with-cross-validation-in-r/>

#Set up cross validation control using 10-fold cross validation. The model will be trained 10 times, each time using a different subset as a test set.

```
train_control <- trainControl(
```

```
  method = "cv", #choose the cross validation method
```

```
  number=10, #choose 10 folds,
```

```
  savePredictions = "final" #parameter ensures that all the predictions are saved
```

```
)
```

#I want to explore all the possible combinations that are made from all the k-nearest-neighbors, the kernels, the distance (Manhattan and Euclidean), I used function expand.grid():

#This idea is referenced from chatGPT:

```
tune_grid <- expand.grid(
  kmax= 1:22, #all values from 1 to 22 for the number of neighbor
  distance=c(1,2), #testing Manhattan distance at 1 and Euclidean distance at 2
  kernel=c("optimal", "triangular", "gaussian") #testing different kernel type in kkn
)
```

```
train_set$R1 <- factor(train_set$R1)
```

#train the k-nearest-neighbor using cross validation:

```
kknmodel <- train(
  R1~., #response
  data = train_set, #training data
  method = "kkn",
  trControl = train_control,
  tuneGrid = tune_grid, #test the combinations in the grid
  preProcess = c("center", "scale") #This code is referenced from chatGPT, to make sure the
  data is centered (subtracted by the mean), and is scaled (divided by standard deviation) before
  training.
)
```

```
print(kknmodel)
```

```
# Extract the results from the train object
results <- kknmodel$results
```

```
# Find the row with the highest accuracy
best_row <- results[which.max(results$Accuracy), ]
```

```
# Print the best hyperparameters and accuracy
print(best_row)
```

```
> print(best_row)
      kmax distance kernel Accuracy      Kappa AccuracySD      KappaSD
115     20         1 optimal 0.8688406 0.7362902 0.04314433 0.08658532
```

```
> print(kknmodel)
k-Nearest Neighbors

457 samples
10 predictor
2 classes: '0', '1'

Pre-processing: centered (10), scaled (10)
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 411, 411, 412, 412, 411, 412, ...
Resampling results across tuning parameters:
```

kmax	distance	kernel	Accuracy	Kappa
1	1	optimal	0.8139130	0.6234456
1	1	triangular	0.8139130	0.6234456
1	1	gaussian	0.8139130	0.6234456
1	2	optimal	0.8183092	0.6327494
1	2	triangular	0.8183092	0.6327494
1	2	gaussian	0.8183092	0.6327494
2	1	optimal	0.8139130	0.6234456
2	1	triangular	0.8139130	0.6234456
2	1	gaussian	0.8139130	0.6234456
2	2	optimal	0.8183092	0.6327494
2	2	triangular	0.8183092	0.6327494
2	2	gaussian	0.8183092	0.6327494
3	1	optimal	0.8139130	0.6234456
3	1	triangular	0.8249275	0.6456788
3	1	gaussian	0.8448309	0.6870707
3	2	optimal	0.8183092	0.6327494
3	2	triangular	0.8335266	0.6638495
3	2	gaussian	0.8475604	0.6833816

So the best k nearest number is 20 with the accuracy of 86.88%